

```
import numpy as np
import pandas as pd
```

```
dataset = pd.read_csv('auto_prize.csv')
dataset.head()
```

	symboling	normalized-losses	wheel-base	length	width
height \					
0	5	164	99.800003	176.600006	66.199997
54.299999					
1	5	164	99.400002	176.600006	66.400002
54.299999					
2	4	158	105.800003	192.699997	71.400002
55.700001					
3	4	158	105.800003	192.699997	71.400002
55.900002					
4	5	192	101.199997	176.800003	64.800003
54.299999					

	curb-weight	engine-size	bore	stroke	compression-ratio
horsepower \					
0	2337	109	3.19	3.4	10.0
102					
1	2824	136	3.19	3.4	8.0
115					
2	2844	136	3.19	3.4	8.5
110					
3	3086	131	3.13	3.4	8.3
140					
4	2395	108	3.50	2.8	8.8
101					

	peak-rpm	city-mpg	highway-mpg	target
0	5500	24	30	13950
1	5500	18	22	17450
2	5500	19	25	17710
3	5500	17	20	23875
4	5800	23	29	16430

```
dataset.describe()
```

	symboling	normalized-losses	wheel-base	length
width \				
count	159.000000	159.000000	159.000000	159.000000
159.000000				
mean	3.735849	121.132075	98.264151	172.413837
65.607547				
std	1.193086	35.651285	5.167417	11.523177

```

1.947883
min      1.000000      65.000000      86.599998      141.100006
60.299999
25%      3.000000      94.000000      94.500000      165.650002
64.000000
50%      4.000000      113.000000      96.900002      172.399994
65.400002
75%      5.000000      148.000000      100.799999      177.800003
66.500000
max      6.000000      256.000000      115.599998      202.600006
71.699997

```

```

count      height      curb-weight      engine-size      bore      stroke \
mean      53.899371      2461.138365      119.226415      3.300126      3.236352
std      2.268761      481.941321      30.460791      0.267336      0.294888
min      49.400002      1488.000000      61.000000      2.540000      2.070000
25%      52.250000      2065.500000      97.000000      3.050000      3.105000
50%      54.099998      2340.000000      110.000000      3.270000      3.270000
75%      55.500000      2809.500000      135.000000      3.560000      3.410000
max      59.799999      4066.000000      258.000000      3.940000      4.170000

```

```

count      compression-ratio      horsepower      peak-rpm      city-mpg
highway-mpg \
count      159.000000      159.000000      159.000000      159.000000
159.000000
mean      10.161132      95.836478      5113.836478      26.522013
32.081761
std      3.889475      30.718583      465.754864      6.097142
6.459189
min      7.000000      48.000000      4150.000000      15.000000
18.000000
25%      8.700000      69.000000      4800.000000      23.000000
28.000000
50%      9.000000      88.000000      5200.000000      26.000000
32.000000
75%      9.400000      114.000000      5500.000000      31.000000
37.000000
max      23.000000      200.000000      6600.000000      49.000000
54.000000

```

```

count      target
mean      11445.729560
std      5877.856195
min      5118.000000
25%      7372.000000
50%      9233.000000
75%      14719.500000
max      35056.000000

```

```
#splitting feature and target cols
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, -1].values
```

Finding correlation between the cols of the dataset

```
dataset.corr()
```

	symboling	normalized-losses	wheel-base	length
\				
symboling	1.000000	0.518344	-0.520591	-0.336257
normalized-losses	0.518344	1.000000	-0.060086	0.035541
wheel-base	-0.520591	-0.060086	1.000000	0.871535
length	-0.336257	0.035541	0.871535	1.000000
width	-0.219186	0.109726	0.814991	0.838338
height	-0.475185	-0.413702	0.555767	0.499251
curb-weight	-0.251880	0.125858	0.810182	0.871291
engine-size	-0.109453	0.207820	0.649206	0.725953
bore	-0.256469	-0.031558	0.578159	0.646318
stroke	-0.021285	0.063330	0.167449	0.121073
compression-ratio	-0.138316	-0.127259	0.291431	0.184814
horsepower	-0.003949	0.290511	0.516948	0.672063
peak-rpm	0.199106	0.237697	-0.289234	-0.234074
city-mpg	0.089550	-0.235523	-0.580657	-0.724544
highway-mpg	0.149830	-0.188564	-0.611750	-0.724599
target	-0.162794	0.202761	0.734419	0.760952

	width	height	curb-weight	engine-size
bore \				
symboling	-0.219186	-0.475185	-0.251880	-0.109453
0.256469				
normalized-losses	0.109726	-0.413702	0.125858	0.207820
0.031558				
wheel-base	0.814991	0.555767	0.810182	0.649206
0.578159				

length 0.646318	0.838338	0.499251	0.871291	0.725953	
width 0.572554	1.000000	0.292706	0.870595	0.779253	
height 0.254836	0.292706	1.000000	0.367052	0.111083	
curb-weight 0.645792	0.870595	0.367052	1.000000	0.888626	
engine-size 0.595737	0.779253	0.111083	0.888626	1.000000	
bore 1.000000	0.572554	0.254836	0.645792	0.595737	
stroke 0.102581	0.196619	-0.091313	0.173844	0.299683	-
compression-ratio 0.015119	0.258752	0.233308	0.224724	0.141097	
horsepower 0.560239	0.681872	0.034317	0.790095	0.812073	
peak-rpm 0.312269	-0.232216	-0.245864	-0.259988	-0.284686	-
city-mpg 0.590440	-0.666684	-0.199738	-0.762155	-0.699139	-
highway-mpg 0.590850	-0.693338	-0.226136	-0.789338	-0.714095	-
target 0.533890	0.843371	0.244836	0.893639	0.841496	
	stroke	compression-ratio	horsepower	peak-	
rpm \ symboling	-0.021285	-0.138316	-0.003949	0.199106	
normalized-losses	0.063330	-0.127259	0.290511	0.237697	
wheel-base	0.167449	0.291431	0.516948	-0.289234	
length	0.121073	0.184814	0.672063	-0.234074	
width	0.196619	0.258752	0.681872	-0.232216	
height	-0.091313	0.233308	0.034317	-0.245864	
curb-weight	0.173844	0.224724	0.790095	-0.259988	
engine-size	0.299683	0.141097	0.812073	-0.284686	
bore	-0.102581	0.015119	0.560239	-0.312269	
stroke	1.000000	0.243587	0.148804	-0.011312	
compression-ratio	0.243587	1.000000	-0.162305	-0.416769	

horsepower	0.148804	-0.162305	1.000000	0.074057
peak-rpm	-0.011312	-0.416769	0.074057	1.000000
city-mpg	-0.020055	0.278332	-0.837214	-0.052929
highway-mpg	-0.012934	0.221483	-0.827941	-0.032777
target	0.160664	0.209361	0.759874	-0.171916

	city-mpg	highway-mpg	target
symboling	0.089550	0.149830	-0.162794
normalized-losses	-0.235523	-0.188564	0.202761
wheel-base	-0.580657	-0.611750	0.734419
length	-0.724544	-0.724599	0.760952
width	-0.666684	-0.693338	0.843371
height	-0.199738	-0.226136	0.244836
curb-weight	-0.762155	-0.789338	0.893639
engine-size	-0.699139	-0.714095	0.841496
bore	-0.590440	-0.590850	0.533890
stroke	-0.020055	-0.012934	0.160664
compression-ratio	0.278332	0.221483	0.209361
horsepower	-0.837214	-0.827941	0.759874
peak-rpm	-0.052929	-0.032777	-0.171916
city-mpg	1.000000	0.971999	-0.692273
highway-mpg	0.971999	1.000000	-0.720090
target	-0.692273	-0.720090	1.000000

the loosely correlated pair of (feature, output) should be dropped in order to apply linear regression model

```
dataset = dataset.drop(["symboling", "normalized-  
losses", "height", "stroke", "compression-ratio", "peak-rpm"], axis=1)  
dataset.corr()
```

	wheel-base	length	width	curb-weight	engine-size
\ wheel-base	1.000000	0.871535	0.814991	0.810182	0.649206
length	0.871535	1.000000	0.838338	0.871291	0.725953
width	0.814991	0.838338	1.000000	0.870595	0.779253
curb-weight	0.810182	0.871291	0.870595	1.000000	0.888626
engine-size	0.649206	0.725953	0.779253	0.888626	1.000000

bore	0.578159	0.646318	0.572554	0.645792	0.595737
horsepower	0.516948	0.672063	0.681872	0.790095	0.812073
city-mpg	-0.580657	-0.724544	-0.666684	-0.762155	-0.699139
highway-mpg	-0.611750	-0.724599	-0.693338	-0.789338	-0.714095
target	0.734419	0.760952	0.843371	0.893639	0.841496

	bore	horsepower	city-mpg	highway-mpg	target
wheel-base	0.578159	0.516948	-0.580657	-0.611750	0.734419
length	0.646318	0.672063	-0.724544	-0.724599	0.760952
width	0.572554	0.681872	-0.666684	-0.693338	0.843371
curb-weight	0.645792	0.790095	-0.762155	-0.789338	0.893639
engine-size	0.595737	0.812073	-0.699139	-0.714095	0.841496
bore	1.000000	0.560239	-0.590440	-0.590850	0.533890
horsepower	0.560239	1.000000	-0.837214	-0.827941	0.759874
city-mpg	-0.590440	-0.837214	1.000000	0.971999	-0.692273
highway-mpg	-0.590850	-0.827941	0.971999	1.000000	-0.720090
target	0.533890	0.759874	-0.692273	-0.720090	1.000000

scaling the features (because the ranges are varying) and splitting training and testing sets

```
from sklearn.preprocessing import MinMaxScaler
# define min max scaler
scaler = MinMaxScaler()
# transform data
X = scaler.fit_transform(X)

#splitting train and test sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =
train_test_split(X,y,test_size=0.30,random_state=125)
```

MODEL TRAINING: linear regression model is used because the target variable is a continuous value => regression and the features are highly correlated with the target (linear relation)

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train,y_train)

LinearRegression()

y_pred = regressor.predict(X_test)
```

MODEL EVALUATION: r-squared is used as a metric to evaluate the model

```
from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
print('r2 score: ', r2)
```

```
r2 score: 0.8804433019303708
```

R-2 of 88% is a good model (>60%)

MODEL INTERPRETATION: the coefficients($b_1 \dots b_n$) of all the features and the intercept(b_0) from the equation : $y = b_0 + b_1x_1 + b_2x_2 + \dots + b_nx_n$

```
regressor.coef_
```

```
array([ -76.37101268,  874.3034376 , 6030.13766262, -4584.71242111,
        7207.75333962, -837.62703782, 9671.68680982, 12731.16586569,
       -2244.62151685, -4894.5030305 , 2617.00906314, 1607.93402346,
        2242.05407738, -5985.52571376, 1394.35434156])
```

```
regressor.intercept_
```

```
4251.92368904624
```

=> from the coef array we can say that: increase in 1 unit in the 1st feature (wheel_base) will decrease the target(auto_price) by 76.37 times, similarly the coefs of other features carry the same meaning

=> the predicted price of the automobile is 4251.92(intercept) units when hypothetically all the features are 0